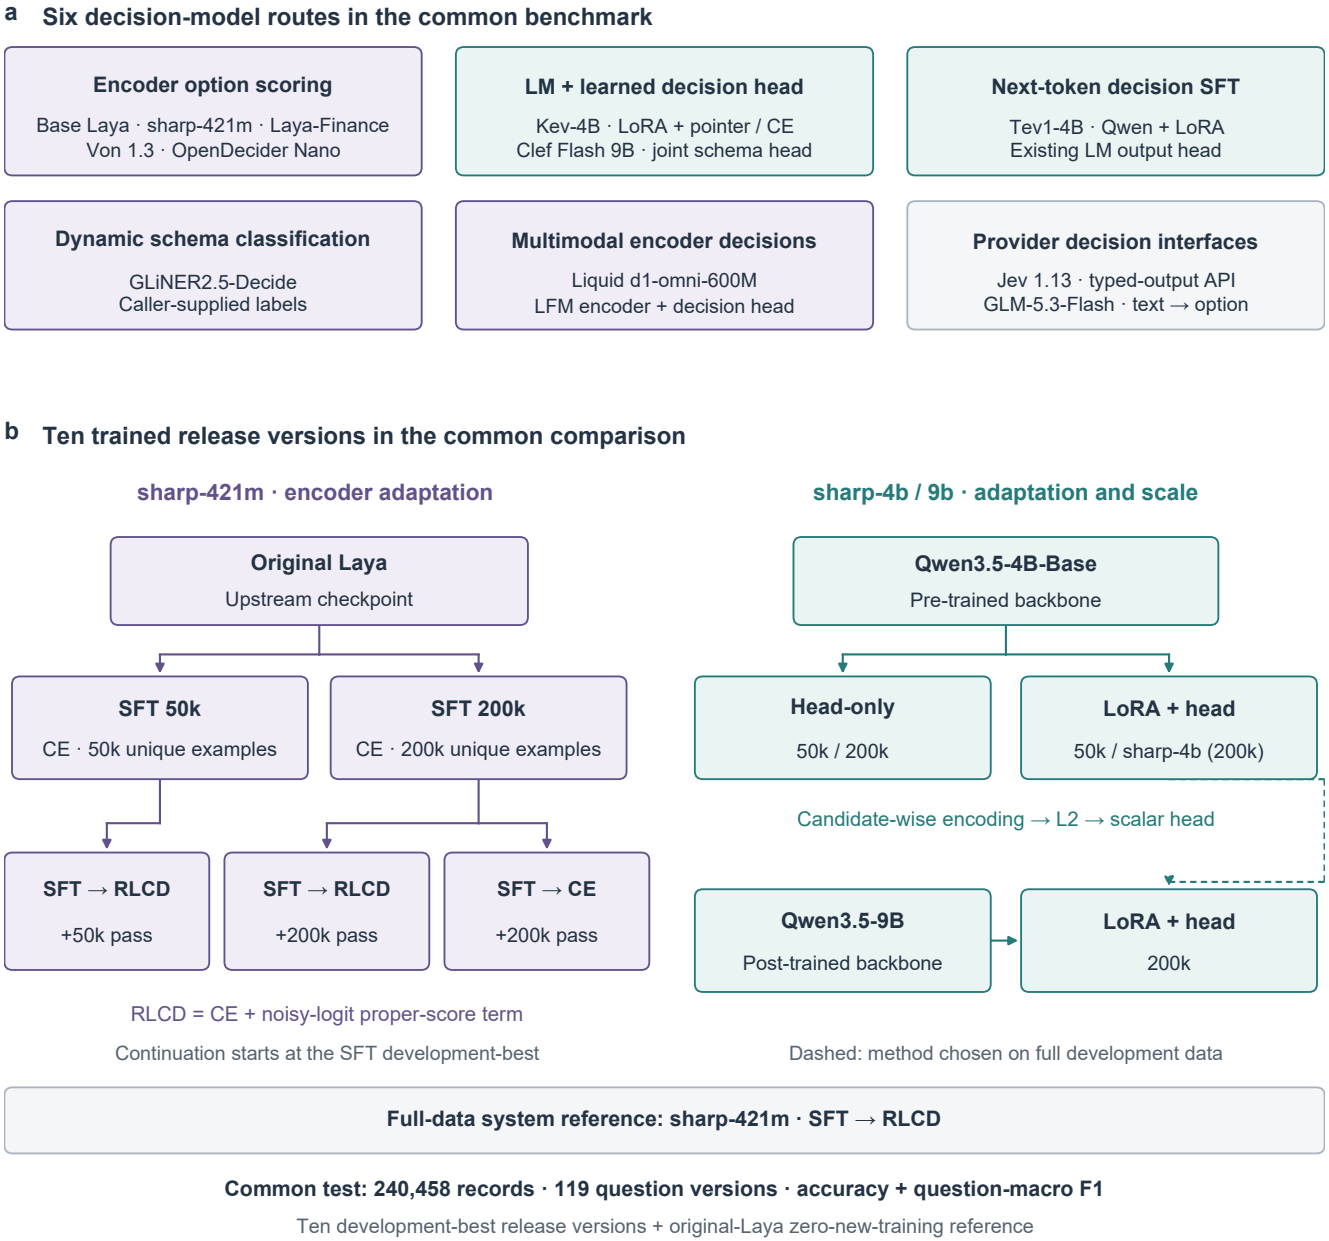

a
Six decision-model routes in the common benchmark
Encoder option scoring
Base Laya · sharp-421m · Laya-Finance
Von 1.3 · OpenDecider Nano
LM + learned decision head
Kev-4B · LoRA + pointer / CE
Clef Flash 9B · joint schema head
Next-token decision SFT
Tev1-4B · Qwen + LoRA
Existing LM output head
Dynamic schema classification
GLiNER2.5-Decide
Caller-supplied labels
Multimodal encoder decisions
Liquid d1-omni-600M
LFM encoder + decision head
Provider decision interfaces
Jev 1.13 · typed-output API
GLM-5.3-Flash · text → option
b
Ten trained release versions in the common comparison
sharp-421m · encoder adaptation
Original Laya
Upstream checkpoint
SFT 50k
CE · 50k unique examples
SFT 200k
CE · 200k unique examples
SFT → RLCD
+50k pass
SFT → RLCD
+200k pass
SFT → CE
+200k pass
RLCD = CE + noisy-logit proper-score term
Continuation starts at the SFT development-best
sharp-4b / 9b · adaptation and scale
Qwen3.5-4B-Base
Pre-trained backbone
Head-only
50k / 200k
LoRA + head
50k / sharp-4b (200k)
Candidate-wise encoding → L2 → scalar head
Qwen3.5-9B
Post-trained backbone
LoRA + head
200k
Dashed: method chosen on full development data
Full-data system reference: sharp-421m · SFT → RLCD
Common test: 240,458 records · 119 question versions · accuracy + question-macro F1
Ten development-best release versions + original-Laya zero-new-training reference